

Data Preprocessing Task

Import Packages

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split
import os
```

Read the Dataset provided (Dataset 1)

- For my own dataset I have used a Air pollution measurement dataset.
- Each rows contain the hourly measurement recording of the 5 pollutants (NO, NO2, NOX, PM10, and PM2.5).
- The data was collected in a location in London for the entirety of year 2017.

```
In [2]: DATA_PATH=os.path.join(os.getcwd(),"LaqnData.csv")
data_1=pd.read_csv(DATA_PATH)
```

```
In [3]: data_1.head() ##preview the dataset using head function
```

```
Out[3]:
```

	Site	Species	ReadingDateTime	Value	Units	Provisional or Ratified
0	H10	CO	01/01/2018 00:00	NaN	mg m-3	P
1	H10	CO	01/01/2018 00:15	NaN	mg m-3	P
2	H10	CO	01/01/2018 00:30	NaN	mg m-3	P
3	H10	CO	01/01/2018 00:45	NaN	mg m-3	P
4	H10	CO	01/01/2018 01:00	NaN	mg m-3	P

```
In [4]: print("shape of dataset ----> ",data_1.shape)
print("Number of examples ----> ",data_1.shape[0])
print("Number of columns ----> ",data_1.shape[1])
```

```
shape of dataset ----> (175200, 6)
Number of examples ----> 175200
Number of columns ----> 6
```

Quick check of the statistics of the dataset using pandas api

As you could see using the describe function it gives us the basic statistics out of it

```
In [5]: data_1.describe()
```

```
Out[5]:
```

	Value
count	131724.000000
mean	48.050503
std	54.252123
min	0.000000
25%	12.500000
50%	33.200000
75%	64.200000
max	643.900020

Info function gives us the information about how many null data we have out of total entries and along with the datatype

```
In [6]: data_1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 175200 entries, 0 to 175199
Data columns (total 6 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Site                                  175200 non-null object
1   Species                              175200 non-null object
2   ReadingDateTime                      175200 non-null object
3   Value                                131724 non-null float64
4   Units                                175200 non-null object
5   Provisional or Ratified              175200 non-null object
dtypes: float64(1), object(5)
memory usage: 8.0+ MB
```

1. Remove Duplicates value

```
In [7]: # shape before removing duplicates
```

```
In [8]: data_1.shape
```

```
Out[8]: (175200, 6)
```

```
In [9]: data_1 = data_1.drop_duplicates()
```

```
In [10]: data_1.shape
```

```
Out[10]: (175200, 6)
```

- So we conclude that there is no duplicates.

2. Remove unwanted columns

- drop function allows to drop any column which is unnecessary and using inplace="True", does not show the column in the dataframe when we see the dataframe back again
- In this case we don't want datetime when it was recorded and the site as it contains single value doesn't provide any insight.

```
In [11]: data_1.drop(columns=['ReadingDateTime', 'Site'], axis=1, inplace=True)
```

3. Checking the null values

```
In [12]: data_1.isna().sum()
```

```
Out[12]: Species          0
Value          43476
Units          0
Provisional or Ratified  0
dtype: int64
```

```
In [13]: ##from the above info function, we can see that Value column has 131724 float values
## even from the above dataframe view we saw Value column has NaN value
## so we need to remove those rows or impute with some impute techniques
## here we would be imputing with mean of the value.
```

```
print('rows present before handling nan :', data_1.shape[0])
data_1['Value'] = data_1['Value'].fillna(data_1['Value'].mean())
print('rows present after handling nan :', data_1.shape[0])
```

```
rows present before handling nan : 175200
rows present after handling nan : 175200
```

4. Encoding of the categorical variable

```
In [14]: data_1['Provisional or Ratified'].value_counts()
```

```
Out[14]: Provisional or Ratified
R      140160
P       35040
Name: count, dtype: int64
```

binary encoding :

- As there are only two values, we will use binary encoding and replace the value with 0 and 1.

```
In [15]: data_1['Provisional or Ratified'] = data_1['Provisional or Ratified'].replac
```

```
In [16]: data_1.head()
```

```
Out[16]:
```

	Species	Value	Units	Provisional or Ratified
0	CO	48.050503	mg m-3	0
1	CO	48.050503	mg m-3	0
2	CO	48.050503	mg m-3	0
3	CO	48.050503	mg m-3	0
4	CO	48.050503	mg m-3	0

```
In [17]: data_1['Species'].value_counts()
```

```
Out[17]: Species
CO      35040
NO      35040
NO2     35040
NOX     35040
O3      35040
Name: count, dtype: int64
```

One-hot encoding:

- As there are multiple value in Species. We will apply one-hot encoding as there is no sequence here.

```
In [18]: data_1 = pd.get_dummies(data_1, columns=['Species'], dtype=int)
```

```
In [19]: data_1.head()
```

```
Out[19]:
```

	Value	Units	Provisional or Ratified	Species_CO	Species_NO	Species_NO2	Species_NOX	Sp
0	48.050503	mg m-3	0	1	0	0	0	
1	48.050503	mg m-3	0	1	0	0	0	
2	48.050503	mg m-3	0	1	0	0	0	
3	48.050503	mg m-3	0	1	0	0	0	
4	48.050503	mg m-3	0	1	0	0	0	

```
In [20]: data_1['Units'].value_counts()
```

```
Out[20]: Units
ug m-3      105120
mg m-3      35040
ug m-3 as N02 35040
Name: count, dtype: int64
```

Again we impliment the same for the Units column

```
In [21]: data_1 = pd.get_dummies(data_1, columns=['Units'], dtype=int)
```

```
In [22]: data_1.head()
```

```
Out[22]:
```

	Value	Provisional or Ratified	Species_CO	Species_NO	Species_NO2	Species_NOX	Species_
0	48.050503	0	1	0	0	0	
1	48.050503	0	1	0	0	0	
2	48.050503	0	1	0	0	0	
3	48.050503	0	1	0	0	0	
4	48.050503	0	1	0	0	0	

5. Split the dataset

```
In [23]: data_train, data_test = train_test_split(data_1, test_size=0.25)
```

show the dataset size after spliting the dataset

```
In [24]: print(data_train.shape)
print(data_test.shape)
```

```
(131400, 10)
(43800, 10)
```

Mean and standard deviation result for numerical columns

- As there is a single column that is numerical we will only calculate the mean and standard deviation for the same column

```
In [25]: Numeric_columns = ['Value']
for col in Numeric_columns:
    print("Column name : ",col)
    print('Training set mean:', data_train[col].mean())
    print('Testing set mean:', data_test[col].mean())
    print('Training set std:', data_train[col].std())
    print('Testing set std:', data_test[col].std())
```

Column name : Value
Training set mean: 48.0201911489353
Testing set mean: 48.141436815864544
Training set std: 46.884353690008936
Testing set std: 47.510488663962995

Conclusion

- We can see that the training datasets mean and standard deviation is slightly less than the test set.

Read the Dataset provided (Dataset 2)

```
In [26]: DATA_PATH_SalaryPrediction=os.path.join(os.getcwd(),"SalaryPrediction.csv")
data_SalaryPrediction=pd.read_csv(DATA_PATH_SalaryPrediction)
```

```
In [27]: data_SalaryPrediction.head() ##preview the dataset using head function
```

```
Out[27]:
```

	Wage	Age	Club	League	Nation	Position	Apps	Caps
0	46,427,000	23.0	PSG	Ligue 1 Uber Eats	FRA	Forward	190	57
1	42,125,000	30.0	PSG	Ligue 1 Uber Eats	BRA	Midfielder	324	119
2	34,821,000	35.0	PSG	Ligue 1 Uber Eats	ARG	Forward	585	162
3	19,959,000	31.0	R. Madrid	La Liga	BEL	Forward	443	120
4	19,500,000	NaN	Man UFC	Premier League	ESP	Goalkeeper	480	45

```
In [28]: data_SalaryPrediction.columns ## columns present in the traffic dataset
```

```
Out[28]: Index(['Wage', 'Age', 'Club', 'League', 'Nation', 'Position', 'Apps', 'Caps'], dtype='object')
```

```
In [29]: print("shape of dataset ----> ",data_SalaryPrediction.shape)
print("Number of examples ----> ",data_SalaryPrediction.shape[0])
print("Number of columns ----> ",data_SalaryPrediction.shape[1])
```

```
shape of dataset ----> (3907, 8)
Number of examples ----> 3907
Number of columns ----> 8
```

Below we would be following same basic techniques to get the statistics as we did in the above

```
In [30]: data_SalaryPrediction.describe()
```

Out [30]:

	Age	Apps	Caps
count	3833.000000	3907.000000	3907.000000
mean	24.060788	140.057077	8.926542
std	4.933452	131.694425	20.518234
min	18.000000	0.000000	0.000000
25%	20.000000	15.000000	0.000000
50%	23.000000	115.000000	0.000000
75%	28.000000	224.500000	6.000000
max	41.000000	715.000000	180.000000

In [31]: `data_SalaryPrediction.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3907 entries, 0 to 3906
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Wage        3907 non-null   object
1    Age         3833 non-null   float64
2    Club        3907 non-null   object
3    League      3907 non-null   object
4    Nation      3907 non-null   object
5    Position    3907 non-null   object
6    Apps        3907 non-null   int64
7    Caps        3907 non-null   int64
dtypes: float64(1), int64(2), object(5)
memory usage: 244.3+ KB
```

1. Remove Duplicates value

In [32]: `# shape before removing duplicates`

In [33]: `data_SalaryPrediction.shape`

Out [33]: (3907, 8)

In [34]: `data_SalaryPrediction = data_SalaryPrediction.drop_duplicates()`

In [35]: `data_SalaryPrediction.shape`

Out [35]: (3842, 8)

- There were some duplicated value which we removed.

2. Remove unwanted columns

- In this dataset, all are important columns. So we will keep all columns for now.

```
In [36]: #data_1.drop(columns=[], axis=1, inplace=True)
```

3. Checking the null values

```
In [37]: data_SalaryPrediction.isna().sum()
```

```
Out[37]: Wage          0
Age          74
Club         0
League       0
Nation       0
Position     0
Apps         0
Caps         0
dtype: int64
```

```
In [ ]:
```

```
In [38]: ##from the above info function, we can see that Value column has 131724 float values
## even from the above dataframe view we saw Value column has NaN value
## so we need to remove those rows or impute with some impute techniques
## here we would be imputing with mean of the value.
```

```
print('rows present before handling nan :', data_SalaryPrediction.shape[0])
data_SalaryPrediction['Age'] = data_SalaryPrediction['Age'].fillna(int(data_SalaryPrediction['Age'].mean()))
print('rows present after handling nan :', data_SalaryPrediction.shape[0])
```

```
rows present before handling nan : 3842
rows present after handling nan : 3842
```

4. Transformation of the data

```
In [39]: data_SalaryPrediction.head()
```

```
Out[39]:
```

	Wage	Age	Club	League	Nation	Position	Apps	Caps
0	46,427,000	23.0	PSG	Ligue 1 Uber Eats	FRA	Forward	190	57
1	42,125,000	30.0	PSG	Ligue 1 Uber Eats	BRA	Midfilder	324	119
2	34,821,000	35.0	PSG	Ligue 1 Uber Eats	ARG	Forward	585	162
3	19,959,000	31.0	R. Madrid	La Liga	BEL	Forward	443	120
4	19,500,000	24.0	Man UFC	Premier League	ESP	Goalkeeper	480	45

- We see that Wage column should be numerical but it is object datatype. So we will remove comma and make it to numerical.

```
In [40]: data_SalaryPrediction['Wage'] = data_SalaryPrediction['Wage'].str.replace(",","")
```

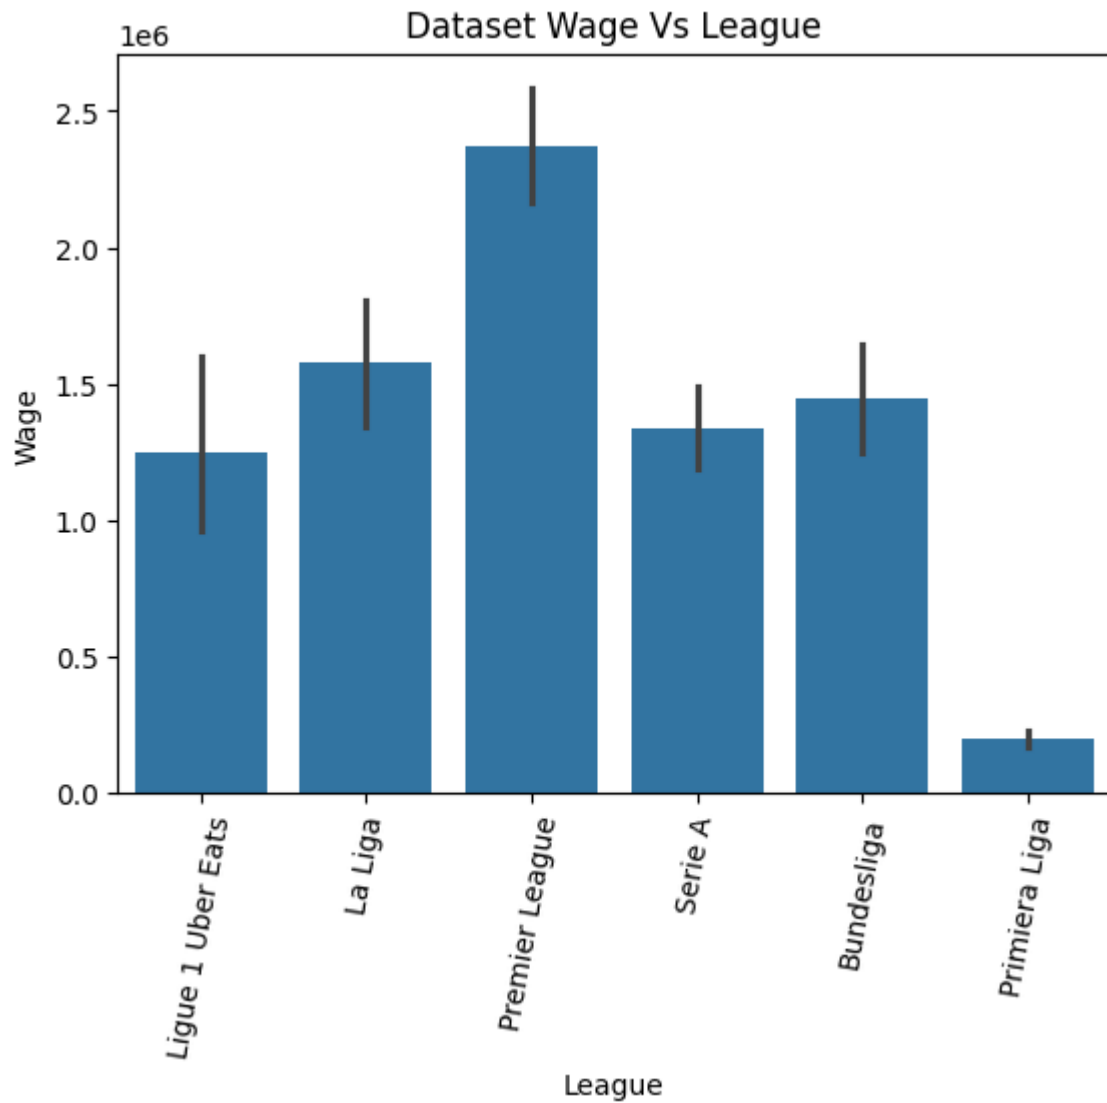

scatter plot wage v/s age

```
In [41]: sns.scatterplot(data=data_SalaryPrediction, x='Age', y='Wage', hue="Wage")
plt.title('Dataset Wage Vs Age')
plt.show()
```



Bar plot wage v/s league

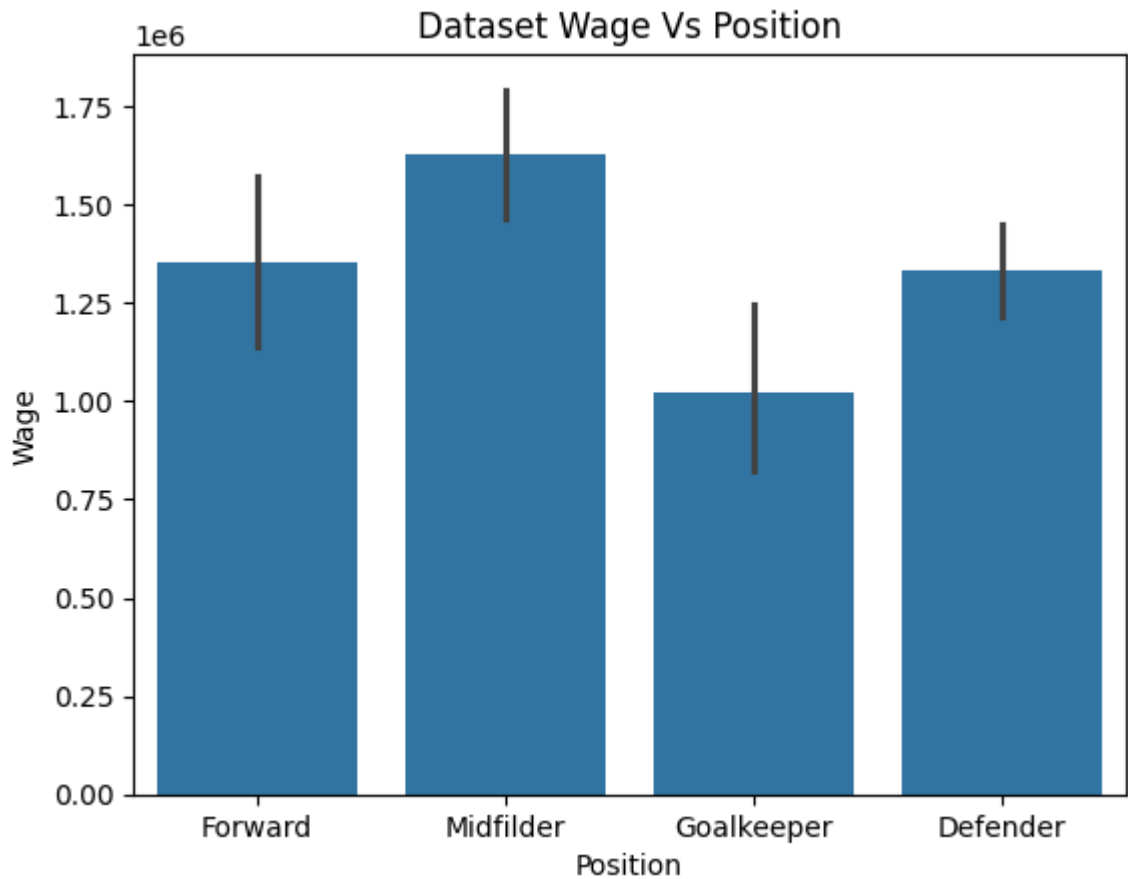
```
In [42]: sns.barplot(data=data_SalaryPrediction, x="League", y="Wage")
plt.title('Dataset Wage Vs League')
plt.xticks(rotation=80)
plt.show()
```



Bar plot wage v/s age

```
In [43]: ## bar plot wage v/s age

sns.barplot(data=data_SalaryPrediction, x="Position", y="Wage")
plt.title('Dataset Wage Vs Position')
plt.show()
```



5. Encoding of the categorical variable

In [44]: `data_SalaryPrediction.head()`

Out[44]:

	Wage	Age	Club	League	Nation	Position	Apps	Caps
0	46427000	23.0	PSG	Ligue 1 Uber Eats	FRA	Forward	190	57
1	42125000	30.0	PSG	Ligue 1 Uber Eats	BRA	Midfielder	324	119
2	34821000	35.0	PSG	Ligue 1 Uber Eats	ARG	Forward	585	162
3	19959000	31.0	R. Madrid	La Liga	BEL	Forward	443	120
4	19500000	24.0	Man UFC	Premier League	ESP	Goalkeeper	480	45

In [45]: `data_SalaryPrediction['Club'].value_counts()`

```
Out[45]: Club
MRT      64
BRG      60
VIZ      55
Chelsea  52
Leicester 51
..
OL       21
FC Lorient 21
LOSC     21
PSG      20
Clermont 18
Name: count, Length: 116, dtype: int64
```

```
In [46]: data_SalaryPrediction['Nation'].value_counts()
```

```
Out[46]: Nation
ESP      431
POR      419
ENG      396
FRA      351
GER      283
...
BDI       1
MSR       1
UZB       1
SUR       1
SIN       1
Name: count, Length: 114, dtype: int64
```

Label Encoding

```
In [47]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Assuming data_SalaryPrediction is your DataFrame
# Creating an instance of LabelEncoder
label_encoder = LabelEncoder()

# Fit and transform the 'Club' column
data_SalaryPrediction['Club'] = label_encoder.fit_transform(data_SalaryPrediction['Club'])
```

```
In [48]: label_encoder = LabelEncoder()

# Fit and transform the 'Club' column
data_SalaryPrediction['Nation'] = label_encoder.fit_transform(data_SalaryPrediction['Nation'])
```

One-hot encoding:

- As there are multiple value but limited in League,Position. We will apply one-hot encoding as there is no sequence here.

```
In [49]: data_SalaryPrediction['League'].value_counts()
```

```
Out[49]: League
Premier League      861
Primiera Liga       736
Serie A              669
La Liga              573
Bundesliga           542
Ligue 1 Uber Eats   461
Name: count, dtype: int64
```

```
In [50]: data_SalaryPrediction['Position'].value_counts()
```

```
Out[50]: Position
Defender      1452
Midfielder    1140
Forward        821
Goalkeeper     429
Name: count, dtype: int64
```

```
In [51]: data_SalaryPrediction = pd.get_dummies(data_SalaryPrediction, columns=['League'])
```

```
In [52]: data_SalaryPrediction.head()
```

```
Out[52]:
```

	Wage	Age	Club	Nation	Apps	Caps	League_Bundesliga	League_La Liga	League_Ligue 1 Uber Eat
0	46427000	23.0	73	40	190	57	0	0	
1	42125000	30.0	73	13	324	119	0	0	
2	34821000	35.0	73	3	585	162	0	0	
3	19959000	31.0	76	8	443	120	0	1	
4	19500000	24.0	64	37	480	45	0	0	

5. Normalising of the data

```
In [53]: data_SalaryPrediction.head()
```

```
Out[53]:
```

	Wage	Age	Club	Nation	Apps	Caps	League_Bundesliga	League_La Liga	League_Ligue 1 Uber Eat
0	46427000	23.0	73	40	190	57	0	0	
1	42125000	30.0	73	13	324	119	0	0	
2	34821000	35.0	73	3	585	162	0	0	
3	19959000	31.0	76	8	443	120	0	1	
4	19500000	24.0	64	37	480	45	0	0	

```
In [54]: numerical_columns = ['Age', 'Apps', 'Caps']
```

```
In [55]: from sklearn.preprocessing import MinMaxScaler
min_maxscaler = MinMaxScaler()
```

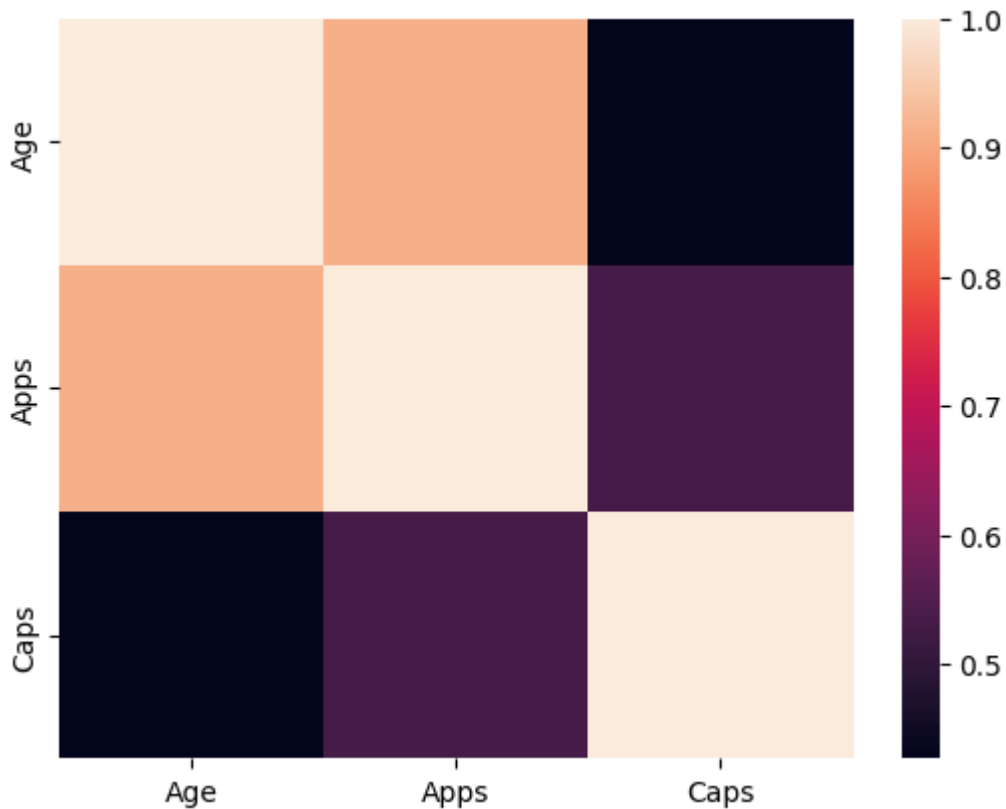
```
In [56]: data_SalaryPrediction[numerical_columns] = min_maxscaler.fit_transform(data
```

6.correlation matrix

to find the correlation between each columns

```
In [57]: ## correlation matrix
data_SalaryPrediction_corr=data_SalaryPrediction[numerical_columns].corr()
sns.heatmap(data_SalaryPrediction_corr)
```

Out[57]: <Axes: >



7. Split the dataset

```
In [58]: data_x = data_SalaryPrediction.drop(['Wage'], axis = 1)
data_y = data_SalaryPrediction['Wage']
data_x_train, data_x_test, data_y_train, data_y_test = train_test_split(data
```

show the dataset size after splitting the dataset

```
In [59]: print(data_x_train.shape, data_y_train.shape)
print(data_x_test.shape, data_y_test.shape)
```

```
(2881, 15) (2881,)
(961, 15) (961,)
```

Mean and standard deviation result

```
In [60]: for col in numerical_columns:
          print(col)
          print('\tTraining set mean:', data_x_train[col].mean())
          print('\tTesting set mean:', data_x_test[col].mean())
          print('\tTraining set std:', data_x_train[col].std())
          print('\tTesting set std:', data_x_test[col].std())
```

Age

```
Training set mean: 0.26517060803163156
Testing set mean: 0.2757091797493553
Training set std: 0.20886794459407704
Testing set std: 0.21938420733236186
```

Apps

```
Training set mean: 0.1973285305461633
Testing set mean: 0.20479977878521063
Training set std: 0.18220882835729854
Testing set std: 0.18907767274145107
```

Caps

```
Training set mean: 0.04928458482779899
Testing set mean: 0.053867499132847725
Training set std: 0.11206559670877406
Testing set std: 0.12250378917178928
```

Conclusion

By this practise, we got to know that data preprocessing is one of crucial steps before putting the data before training. the dataset which we recieve are usually quite noisy, so they required proper preprocessing techniques, such that the model doesn't learn the unwanted information. Putting a lot of time is essential, as cleaning makes the data more readable and model friendly, which ultimately helps us in achiving our goals