

```
In [1]: import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
```

Own dataset

- The dataset contain year wise co2 and temperature difference data.

```
In [2]: df = pd.read_csv("co2_temperature.csv")
df.head()
```

```
Out[2]:
```

	Year	CO2	Temperature_change
0	1961	2580	0.0818
1	1962	2686	0.0924
2	1963	2833	0.1100
3	1964	2995	-0.1461
4	1965	3130	-0.0752

```
In [3]: df.shape
```

```
Out[3]: (50, 3)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Year                  50 non-null    int64
1   CO2                   50 non-null    int64
2   Temperature_change    50 non-null    float64
dtypes: float64(1), int64(2)
memory usage: 1.3 KB
```


```
In [5]: df.describe()
```

Out [5]:

	Year	CO2	Temperature_change
count	50.00000	50.000000	50.000000
mean	1985.50000	5625.820000	0.278896
std	14.57738	1703.260675	0.240433
min	1961.00000	2580.000000	-0.146100
25%	1973.25000	4600.500000	0.092725
50%	1985.50000	5523.000000	0.275400
75%	1997.75000	6634.750000	0.455150
max	2010.00000	9167.000000	0.700800

In [6]:

```
from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure()
fig.set_size_inches(12.5, 9.5)
ax = fig.add_subplot(111, projection='3d')
ax.scatter(xs= df['Year'],ys=df['Temperature_change'],zs = df['CO2']) # scat
ax.set_ylabel('Relative tempature'); ax.set_xlabel('Year'); ax.set_zlabel('C
ax.view_init(10, -45)
plt.show()
```




```
In [7]: fig = plt.figure()
fig.set_size_inches(12.5, 9.5)

plt.subplot(2, 1, 1)
plt.plot(df['Year'],df['CO2'])
plt.ylabel('CO2 Emissions (million metric tons)')

plt.subplot(2, 1, 2)
plt.plot(df['Year'],df['Temperature_change'])
plt.xlabel('Year')
plt.ylabel("Relative temperature change(in 'C)")

plt.show()
```



```
In [8]: df.head()
```

```
Out[8]:
```

	Year	CO2	Temperature_change
0	1961	2580	0.0818
1	1962	2686	0.0924
2	1963	2833	0.1100
3	1964	2995	-0.1461
4	1965	3130	-0.0752

Split the dataset

```
In [9]: X = df['Year'].values
        y = df['Temperature_change'].values
```

```
In [10]: X_train,X_test,y_train,y_test = train_test_split(X.reshape(-1,1), y.reshape(-1,1),
```

Linear Regression

```
In [11]: reg = LinearRegression()
```

```
In [12]: reg.fit(X_train, y_train)
```

```
y_pred = reg.predict(X_test)
```

```
In [13]: MSE = mean_squared_error(y_test, y_pred)
print(f"Mean squared error : {MSE}")
```

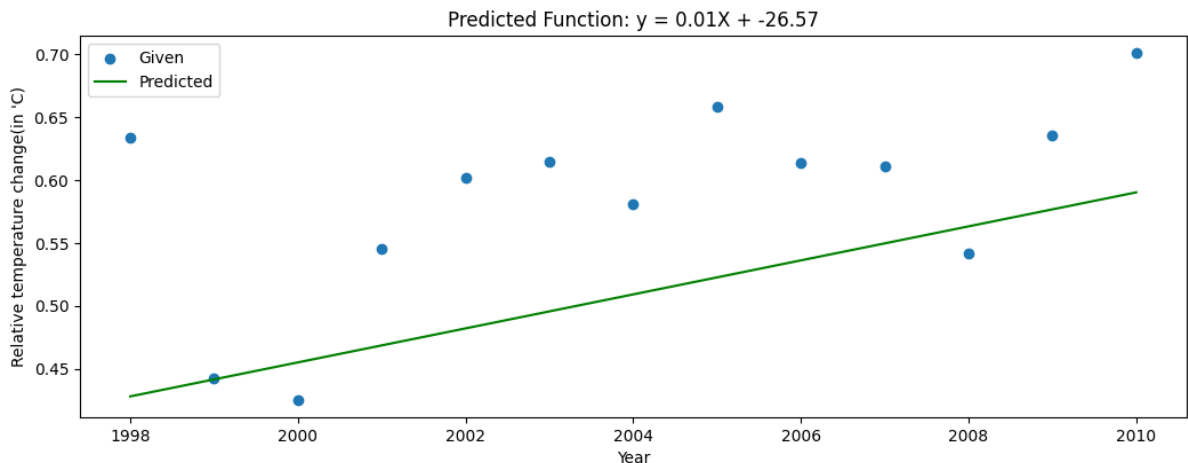
Mean squared error : 0.009763461396018537

```
In [14]: r2_score_ = r2_score(y_test, y_pred)
print(f"R2 score : {r2_score_}")
```

R2 score : -0.6807757772371528

```
In [15]: fig = plt.figure()
fig.set_size_inches(12.5, 9.5)

plt.subplot(2, 1, 1)
plt.scatter(X_test, y_test, label='Given')
plt.plot(X_test, y_pred, color='green', label='Predicted')
plt.xlabel('Year')
plt.ylabel("Relative temperature change(in 'C)")
titlestr = 'Predicted Function: y = %.2fX + %.2f' % (reg.coef_[0], reg.intercept_)
plt.title(titlestr)
plt.legend()
plt.show()
```



Multilinear regression

```
In [16]: X = df.drop(['Temperature_change'], axis=1)
y = df['Temperature_change'].values
```

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(X, y.reshape(-1,1), test_si
```

```
In [18]: # Initialize the standardizer
scaler = StandardScaler()

# Fit on training set only
scaler.fit(X_train)

# Apply transform to both the training set and the test set
```

```
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [19]: reg = LinearRegression()
reg.fit(X_train, y_train)
y_pred = reg.predict(X_test)
```

```
In [20]: MSE = mean_squared_error(y_test, y_pred)
print(f"Mean squared error : {MSE}")

Mean squared error : 0.012639292655687593
```

```
In [21]: r2_score_ = r2_score(y_test, y_pred)
print(f"R2 score : {r2_score_}")


R2 score : -1.1758489203177684
```

```
In [22]: fig = plt.figure()
fig.set_size_inches(12.5, 9.5)

plt.subplot(2, 1, 1)
plt.scatter(X_test[:,0], y_test, label='Given')
plt.plot(X_test[:,0], y_pred, color='green', label='Predicted')
plt.ylabel('CO2 Emissions (million metric tons)')
plt.legend()

plt.subplot(2, 1, 2)
plt.scatter(X_test[:,1], y_test, label='Given')
plt.plot(X_test[:,1], y_pred, color='green', label='Predicted')
plt.xlabel('CO2 Emissions (million metric tons)')
plt.ylabel("Relative temperature change(in 'C)")
plt.legend()

plt.show()
```



Given Dataset

```
In [23]: df = pd.read_csv("Admission_Predict_Ver1.1_small_data_set_for_Linear_Regress
```

```
In [24]: df.head()
```

Out[24]:

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

```
In [25]: df.shape
```

```
Out[25]: (500, 9)
```

```
In [26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Serial No.            500 non-null   int64
1   GRE Score              500 non-null   int64
2   TOEFL Score            500 non-null   int64
3   University Rating      500 non-null   int64
4   SOP                    500 non-null   float64
5   LOR                    500 non-null   float64
6   CGPA                   500 non-null   float64
7   Research               500 non-null   int64
8   Chance of Admit        500 non-null   float64
dtypes: float64(4), int64(5)
memory usage: 35.3 KB
```

```
In [27]: df.describe()
```

```
Out[27]:
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CG
count	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000	500.000000
mean	250.500000	316.472000	107.192000	3.114000	3.374000	3.484000	8.576400
std	144.481833	11.295148	6.081868	1.143512	0.991004	0.925450	0.604000
min	1.000000	290.000000	92.000000	1.000000	1.000000	1.000000	6.800000
25%	125.750000	308.000000	103.000000	2.000000	2.500000	3.000000	8.127500
50%	250.500000	317.000000	107.000000	3.000000	3.500000	3.500000	8.560000
75%	375.250000	325.000000	112.000000	4.000000	4.000000	4.000000	9.040000
max	500.000000	340.000000	120.000000	5.000000	5.000000	5.000000	9.920000

```
In [28]: df = df.drop(['Serial No.'], axis=1)
```

```
In [29]: X = df.drop(['Chance of Admit '],axis=1)
y = df['Chance of Admit ']
```

```
In [30]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

print(f"Number of training examples: {X_train.shape[0]}")
print(f"Number of testing examples: {X_test.shape[0]}")

Number of training examples: 400
Number of testing examples: 100
```

Linear Regression

```
In [31]: reg = LinearRegression()
```

```
In [32]: reg.fit(X_train['CGPA'].values.reshape(-1,1), y_train)
```




```
Out[32]: ▼ LinearRegression
LinearRegression()
```

```
In [33]: y_pred = reg.predict(X_test['CGPA'].values.reshape(-1,1))
```

```
In [34]: reg.intercept_
```

```
Out[34]: -1.0182042839120313
```

```
In [35]: plt.scatter(X_test['CGPA'], y_test, color='black',s=2)
plt.xlabel('CGPA')
plt.ylabel('Chance of Admit ')
plt.plot(X_test['CGPA'], y_pred, color='red')
titlestr = 'Predicted Function: y = %.2fX + %.2f' % (reg.coef_[0], reg.intercept_)
plt.title(titlestr)
plt.show()
```



```
In [36]: MSE = mean_squared_error(y_test, y_pred)
print(f"Mean squared error : {MSE}")
```

Mean squared error : 0.00387225142355091

```
In [37]: r2_score_ = r2_score(y_test, y_pred)
print(f"R2 score : {r2_score_}")
```

R2 score : 0.8290680477081549

Multiple Linear Regression

```
In [38]: X = df.drop(['Chance of Admit '],axis=1)
y = df['Chance of Admit ']
```

```
In [39]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

print(f"Number of training examples: {X_train.shape[0]}")
print(f"Number of testing examples: {X_test.shape[0]}")

Number of training examples: 400
Number of testing examples: 100
```

```
In [40]: # Initialize the standardizer
scaler = StandardScaler()

# Fit on training set only
scaler.fit(X_train)

# Apply transform to both the training set and the test set
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [41]: reg = LinearRegression()
```

```
In [42]: reg.fit(X_train, y_train)
```


```
Out[42]: ▼ LinearRegression
LinearRegression()
```

```
In [43]: y_pred = reg.predict(X_test)
```

```
In [44]: X_test[:, -2].shape
```

```
Out[44]: (100,)
```

```
In [45]: plt.scatter(y_test, y_pred, color='black',s=2)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red')
plt.xlabel('Actual')
plt.ylabel('Perdicted')
plt.show()
```



```
In [46]: MSE = mean_squared_error(y_test, y_pred)
print(f"Mean squared error : {MSE}")
```

Mean squared error : 0.0030965712093589813

```
In [47]: r2_score_ = r2_score(y_test, y_pred)
print(f"R2 score : {r2_score_}")
```

R2 score : 0.8633087306761005

Classification

```
In [48]: df = pd.read_csv("Admission_Predict_Ver1.1_small_data_set_for_Linear_Regress")
df.head()
```

```
Out[48]:
```

	Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	1	337	118	4	4.5	4.5	9.65	1	0.92
1	2	324	107	4	4.0	4.5	8.87	1	0.76
2	3	316	104	3	3.0	3.5	8.00	1	0.72
3	4	322	110	3	3.5	2.5	8.67	1	0.80
4	5	314	103	2	2.0	3.0	8.21	0	0.65

```
In [49]: df.isna().sum()
```

```
Out[49]: Serial No.          0
GRE Score          0
TOEFL Score        0
University Rating  0
SOP                0
LOR                0
CGPA               0
Research           0
Chance of Admit    0
dtype: int64
```

```
In [50]: df = df.drop(['Serial No.'], axis=1)
```

```
In [51]: df['University Rating'].value_counts()
```

```
Out[51]: University Rating
3      162
2      126
4      105
5       73
1       34
Name: count, dtype: int64
```

```
In [52]: df.columns
```

```
Out[52]: Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGP
A',
               'Research', 'Chance of Admit '],
              dtype='object')
```

```
In [53]: # columns_for_transformation = ['GRE Score', 'TOEFL Score', 'SOP', 'LOR ', '
# # Initialize the MinMaxScaler
# scaler = MinMaxScaler()


# # Apply the scaler to the numerical columns
# df[columns_for_transformation] = scaler.fit_transform(df[columns_for_transformation])
```

```
In [54]: df.head()
```

```
Out[54]:
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
0	337	118	4	4.5	4.5	9.65	1	0.92
1	324	107	4	4.0	4.5	8.87	1	0.76
2	316	104	3	3.0	3.5	8.00	1	0.72
3	322	110	3	3.5	2.5	8.67	1	0.80
4	314	103	2	2.0	3.0	8.21	0	0.65

```
In [55]: df['Chance of Admit '].hist(bins=10)
plt.title("Histogram of Chances of admit")
plt.show()
```



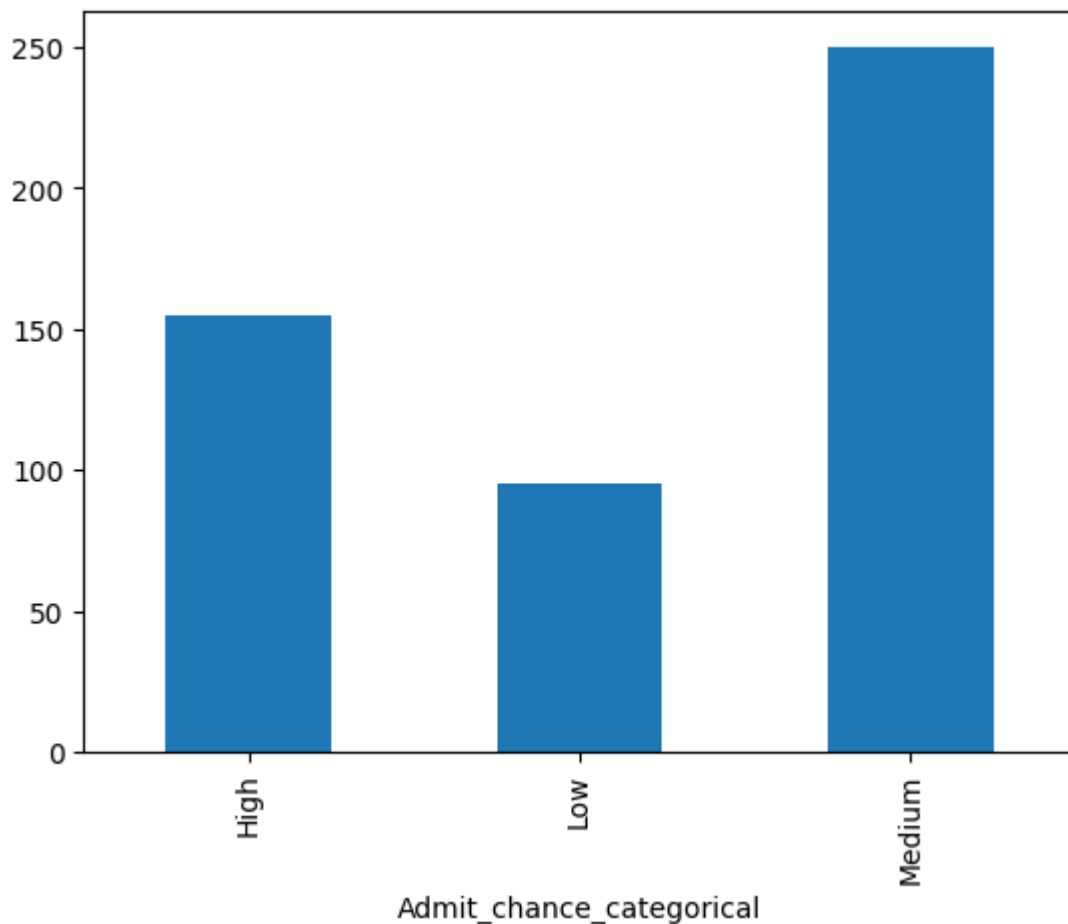
```
In [56]: def convert_to_categorical(chance_probability):
            if chance_probability >= 0.8:
                return "High"
            elif chance_probability >= 0.6:
                return "Medium"
            else:
                return "Low"

            df['Admit_chance_categorical'] = df['Chance of Admit '].apply(convert_to_cat
            df.drop(['Chance of Admit '],axis=1, inplace=True)
            df.head()
```

```
Out[56]:
```

	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Admit_chance_categorical
0	337	118	4	4.5	4.5	9.65	1	High
1	324	107	4	4.0	4.5	8.87	1	Medium
2	316	104	3	3.0	3.5	8.00	1	Medium
3	322	110	3	3.5	2.5	8.67	1	High
4	314	103	2	2.0	3.0	8.21	0	Medium

```
In [57]: df['Admit_chance_categorical'].value_counts().sort_index().plot(kind='bar')
plt.show()
```



Split the dataset

```
In [58]: X = df.drop(['Admit_chance_categorical'],axis=1)
y = df['Admit_chance_categorical']
```

```
In [59]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran
print(f"Number of training examples: {X_train.shape[0]}")
print(f"Number of testing examples: {X_test.shape[0]}")
```

Number of training examples: 400
Number of testing examples: 100

```
In [60]: clf = DecisionTreeClassifier(max_depth=3, random_state = 0)
clf.fit(X_train, y_train)
```


```
Out[60]: ▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3, random_state=0)
```

```
In [61]: y_pred = clf.predict(X_test)
```

```
In [62]: accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy : {accuracy}")
```

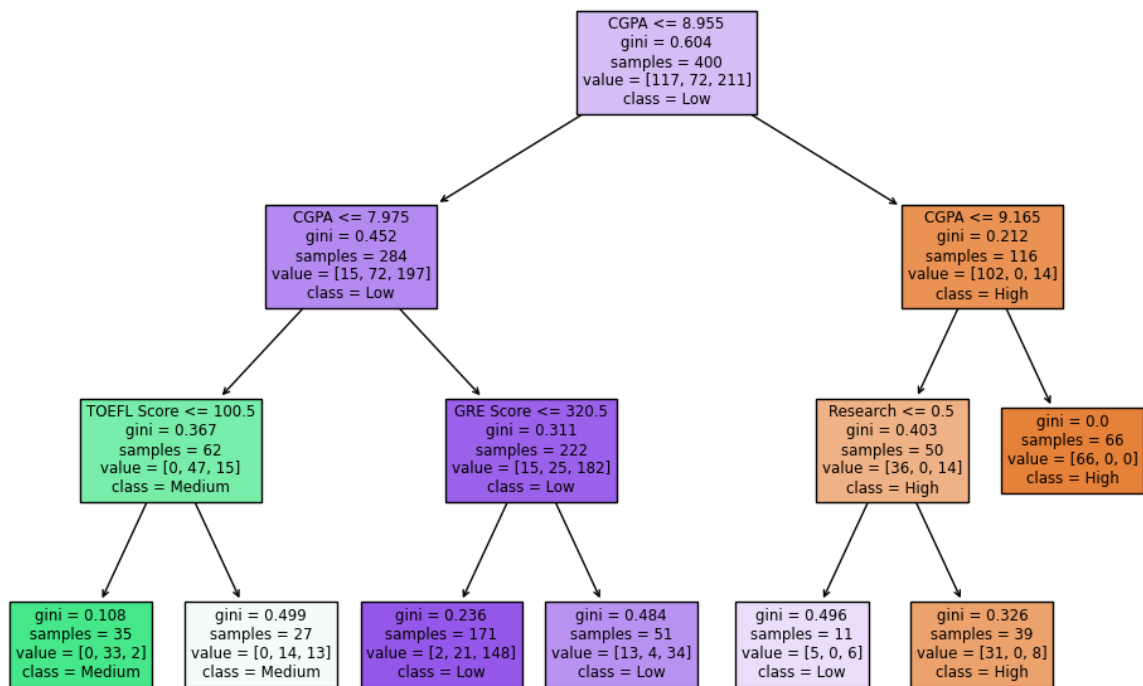
Accuracy : 0.78

```
In [63]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.show()
```



```
In [64]: import matplotlib.pyplot as plt
from sklearn.tree import plot_tree

# Plot the decision tree
plt.figure(figsize=(12, 8)) # Set the size of the figure
plot_tree(clf, filled=True, feature_names=X_test.columns, class_names=["High", "Low", "Medium"])
plt.show()
```



In [65]: `from sklearn.tree import export_text`

```
tree_rules = export_text(clf, feature_names=list(X_test.columns))
print(tree_rules)
```

```

|--- CGPA <= 8.95
|   |--- CGPA <= 7.97
|   |   |--- TOEFL Score <= 100.50
|   |   |   |--- class: Low
|   |   |   |--- TOEFL Score > 100.50
|   |   |   |   |--- class: Low
|   |   |--- CGPA > 7.97
|   |   |   |--- GRE Score <= 320.50
|   |   |   |   |--- class: Medium
|   |   |   |   |--- GRE Score > 320.50
|   |   |   |   |--- class: Medium
|   |--- CGPA > 8.95
|   |   |--- CGPA <= 9.16
|   |   |   |--- Research <= 0.50
|   |   |   |   |--- class: Medium
|   |   |   |   |--- Research > 0.50
|   |   |   |   |--- class: High
|   |   |--- CGPA > 9.16
|   |   |   |--- class: High

```

Conclusion

- From the above rule it is pretty much clear that above CGPA 9.16, the chance of admission is high.

- If the CGPA is above 8.95 and less than 9.16 and someone has an experience in research then also the chance of admission is high.
- If CGPA score more than 7.97 and less than 8.95 then the chances are medium.
- Similarly if CGPA is less than 7.97, than chances are Low.

C: Classification Tree Part

```
In [66]: import pandas as pd
import numpy as np
```

- We copy the data from pdf to a csv file.

```
In [67]: df = pd.read_csv("task3.csv")
df
```

```
Out[67]:
```

	color	shape	size	class
0	red	square	big	+
1	blue	square	big	+
2	red	round	small	-
3	green	square	small	-
4	red	round	big	+
5	green	round	big	-

Calculating Initial Entropy

```
In [68]: def calculate_intial_entropy(dataset, target_column, unique_classes):
total_entries = dataset[target_column].shape[0]
class_counts = dataset[target_column].value_counts()
entropy = -(class_counts / total_entries) * np.log2((total_entries - cla
total_entropy = entropy.sum()
return total_entropy
```

```
In [69]: class_entropy = calculate_intial_entropy(df, 'class', ['+', '-']) #Class
print(class_entropy)
```

1.0

entropy and information gain

```
In [70]: import numpy as np
import pandas as pd

def calculate_average_entropy(dataset, target_attr, unique_classes):
total_entries = dataset.shape[0]
average_entropy = 0
```

```

for class_value in unique_classes:
    class_subset = dataset[dataset[target_attr] == class_value]
    class_count = class_subset.shape[0]

    if class_count == 0:
        continue

    positive_count = class_subset[class_subset['class'] == '+'].shape[0]
    negative_count = class_count - positive_count

    # Calculate entropy for both positive and negative instances
    positive_entropy = -(positive_count / class_count) * np.log2(positive_count / class_count)
    negative_entropy = -(negative_count / class_count) * np.log2(negative_count / class_count)

    # Total entropy for the class
    class_entropy = positive_entropy + negative_entropy

    # Weighted average entropy
    average_entropy += (class_count / total_entries) * class_entropy

return round(average_entropy, 2)

```

```

In [71]: color_entropy = calculate_average_entropy(df, 'color', ['red', 'green', 'blue'])
print("Color-> entropy: ", color_entropy, ", information gain: ", 1 - color_entropy)
shape_entropy = calculate_average_entropy(df, 'shape', ['square', 'round'])
print("Shape-> entropy: ", shape_entropy, ", information gain: ", round(1 - shape_entropy, 2))
size_entropy = calculate_average_entropy(df, 'size', ['big', 'small'])
print("Size-> entropy: ", size_entropy, ", information gain: ", round(1 - size_entropy, 2))

Color-> entropy: 0.46 , information gain: 0.54
Shape-> entropy: 0.92 , information gain: 0.08
Size-> entropy: 0.54 , information gain: 0.46

```

Decision Tree

```

In [72]: from sklearn.preprocessing import LabelEncoder
import matplotlib.pyplot as plt
from sklearn.tree import plot_tree

le = LabelEncoder()

for column in df.columns:
    df[column] = le.fit_transform(df[column])

df

```

Out [72]:

	color	shape	size	class
0	2	1	0	0
1	0	1	0	0
2	2	0	1	1
3	1	1	1	1
4	2	0	0	0
5	1	0	0	1

```
In [73]: X = df[['color', 'shape', 'size']]  
y = df['class']
```

```
In [74]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.25, ran
```


```
In [75]: clf_entropy = DecisionTreeClassifier(random_state=10)  
clf_entropy.fit(X_train, y_train)
```

```
Out[75]: ▼ DecisionTreeClassifier  
DecisionTreeClassifier(random_state=10)
```

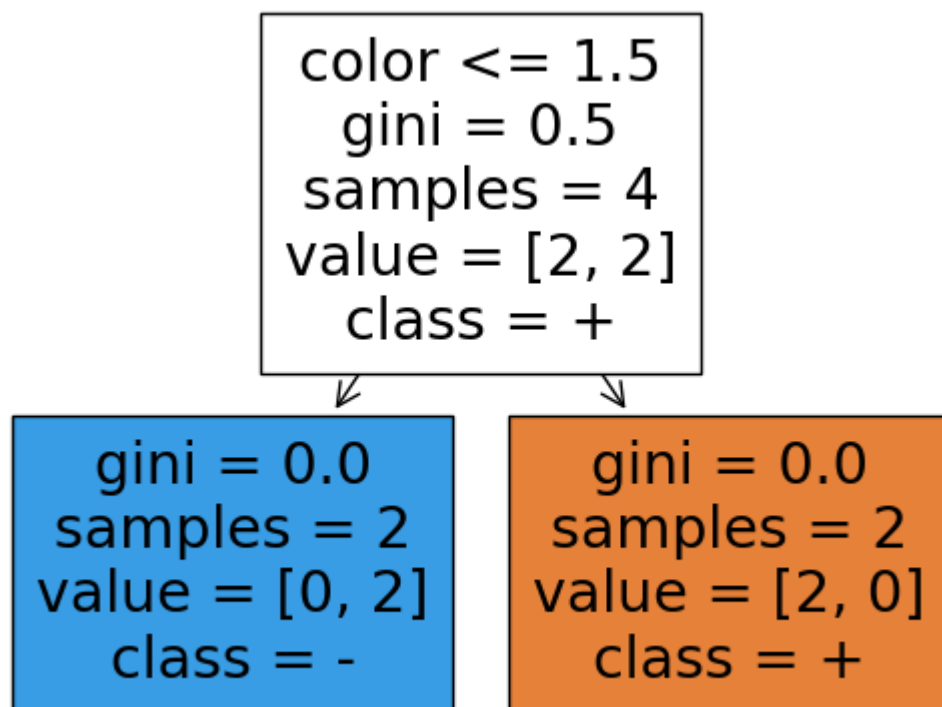
```
In [76]: y_pred = clf_entropy.predict(X_test)
```

```
In [77]: from sklearn.metrics import ConfusionMatrixDisplay  
print("Accuracy score on test: ", clf_entropy.score(X_test, y_test))  
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)  
plt.show()
```

Accuracy score on test: 0.0



```
In [78]: plot_tree(clf_entropy, filled=True, feature_names=X_test.columns, class_name=
plt.show())
```



```
In [79]: from sklearn.tree import export_text
```

```
tree_rules = export_text(clf_entropy, feature_names=list(X_test.columns))
print(tree_rules)
```

```
|--- color <= 1.50
|   |--- class: 1
|--- color > 1.50
|   |--- class: 0
```

Conclusion

- As we can clearly see that the color is splitted to make the tree as it has the highest information gain.

In []: